

WASHINGTON UNIVERSITY IN ST. LOUIS
School of Engineering and Applied Science
Department of Computer Science and Engineering

Masters Project Examination Committee:

Tao Ju
Robert Pless
Bill Smart

TOOLS FOR VISUALIZING WEBCAM SCENES

by

David Ross

A masters project presented to the School of Engineering
of Washington University in partial fulfillment of the
requirements for the degree of

MASTER OF SCIENCE

May 2009
Saint Louis, Missouri

ABSTRACT OF THE THESIS

Tools for Visualizing Webcam Scenes

by

David Ross

Master of Science in Computer Science

Washington University in St. Louis, 2009

Research Advisor: Professor Robert Pless

Webcams are pretty awesome. Here's how to visualize them...

Acknowledgments

First, I want to thank my adviser, Professor Robert Pless for four years of teaching, guidance, and support. Without his direction, this project would not have been completed, and my four years here would not have been as successful.

I want to thank the graduate students in the Media and Machines Lab, especially Nathan Jacobs and Michael Dixon, for helping me along the way. Very rarely did I have a question that went unanswered for long. Also, thanks to Tao Ju and Bill Smart for agreeing to be on my project committee. I hope you enjoyed it.

David Ross

Washington University in Saint Louis
May 2009

Contents

Abstract	ii
Acknowledgments	iii
List of Figures	vi
1 Introduction	1
1.1 Related Work	1
1.2 AMOS Dataset	1
2 Principal Component Analysis	2
2.1 Background	2
2.2 Theory	2
2.3 Algorithm	2
2.3.1 Incremental PCA	2
2.4 Applications	2
3 Visualization Tools	3
3.1 Setup	3
3.1.1 Temporal Narrowing	3
3.1.2 Sky Mask	4
3.1.3 Gradient Image	5
3.2 Visualizations	6
3.2.1 Image Montage	6
3.2.2 Intelligent Image Montage	6
3.2.3 Two-dimensional Explorer	8
3.3 Criteria	8
3.3.1 PCA coefficient vector magnitude	9
3.3.2 Residual Error	9
3.3.3 Variance Model	9
3.3.4 Probabilistic	14
4 Discussion	17
4.1 Future Work	17
References	18

List of Figures

3.1	Learning a sky mask for a webcam scene.	4
3.2	Focusing on object edges with gradient magnitude images.	5
3.3	Image montage techniques.	7
3.4	Exploring 2D image space using a simple GUI.	8
3.5	Criteria: Gradient Magnitude.	10
3.6	Residual Reconstruction.	11
3.7	Residual SSD Montage.	11
3.8	Several variance images.	12
3.9	Z-Score Montage.	13
3.10	Several residual image histograms.	14
3.11	Least Gaussian and Least Laplacian Images.	15

Chapter 1

Introduction

Webcams are cool

1.1 Related Work

1.2 AMOS Dataset

Chapter 2

Principal Component Analysis

Principal Component Analysis (PCA) is a powerful tool for

2.1 Background

11

2.2 Theory

2.3 Algorithm

2.3.1 Incremental PCA

2.4 Applications

Chapter 3

Visualization Tools

The main contribution of this project is automatic methods for visualizing how a webcam scene varies, and for understanding the most important variations. In most natural scenes, we notice changes in lighting, weather, and camera conditions which are interesting, but fail to describe the typical behavior of a scene. The goal of these tools is to learn these variations and to point out changes independent of them. PCA is a commonly used tool, and captures a linear model of consistent image variations. By analyzing the results of certain PCA setups, we can obtain important and interesting information.

3.1 Setup

The most obvious way to learn about scenes is to take the PCA decomposition of the entire webcam scene. Unfortunately, our data-set is too large for this to be feasible, so we must limit our inputs. Instead of taking a random subset of the images, if we can intelligently limit the images we give, we can get better inputs.

3.1.1 Temporal Narrowing

The AMOS data-set consists mostly of outdoor scenes. These outdoor scenes vary significantly over the course of a day, going from night to day and back. The change in lighting dominates all other changes across the scene, and causes many images to be completely dark. Instead of wasting

3.1.2 Sky Mask

In many outdoor scenes, even when narrowed to a particular time of day, the most difficult image variation to characterize is the sky. PCA has a difficult time learning changes in sunlight, clouds, and other characteristics of the sky, even among a set of images of one time of day. The variance in the sky dominates the PCA reconstruction, causing it to ignore information that is more interesting in this context.

Fortunately, it is fairly easy to learn which regions of an image are most affected by this. In practice, a side effect of the rising and setting of the sun is that the first principal component of most scenes is the sky, as shown in 3.1(a). By thresholding the values of this vector, we can effectively segment the sky from the rest of the image, as shown in 3.1(b).

This simple mask allows PCA to focus on more interesting changes in the image, and most of the results that follow use a simple mask to ignore the sky regions of the images. By setting all pixels in the sky to 0, we can effectively remove this difficulty.



Figure 3.1: Figure 3.1(a) shows the first PCA component of a webcam scene. By adding or subtracting this component, we control how dark or how light the sky is. Figure 3.1(b) shows that a simple thresholding of this image effectively segments the sky from the rest of the image.

3.1.3 Gradient Image

Webcam images are very high-dimensional - a typical 320 x 240 gray scale image has 76,800 pixels, each of which is a value from 0 to 255. One way to simplify this space without changing the size of the image is to look at the gradient magnitude images of a scene.

The x and y derivatives of an image are defined as $I_x(x, y) = I(x, y) - I(x - 1, y)$ and $I_y(x, y) = I(x, y) - I(x, y - 1)$ where the function $I(x, y)$ describes the intensities of an image's pixels. Once we have calculated the derivative images, we define the gradient magnitude of an image to be $G(x, y) = I_x(x, y)^2 + I_y(x, y)^2$.

The gradient magnitude image highlights edges in an image, as they are image locations where pixel values differ greatly from their neighbors. By performing PCA on the gradient magnitude images, we tend to ignore the potentially noisy surfaces of objects, and instead focus on the locations of the objects. Figure 3.2(a) and Figure 3.2(b) show the edges of an interesting webcam image highlighted in a gradient magnitude image.



(a)



(b)

Figure 3.2: Figure 3.2(a) shows a grayscale webcam image. Figure 3.2(b) shows the gradient magnitude image of that frame. Notice the noise is mostly removed and the edges are highlighted.

3.2 Visualizations

In this section, we present several tools to visualize the webcam scene based on criteria we will discuss in the next section.

3.2.1 Image Montage

The simplest way to visualize a webcam scene is to view a montage of several images from that scene. We can easily sort the images along some dimension, and then show the n images with the highest values in that dimension. Figure 3.3(a) Is an example montage of several interesting images from a webcam scene.

3.2.2 Intelligent Image Montage

In several of these montages, we notice that many of the most unusual images are similar to each other. This presents the problem of finding unusual images that are different from the images we have already chosen.

A simple but effective algorithm for this problem uses the L2 norm in image space. If we consider an image as a vector of pixel intensities, such as $v = \{v_1, \dots, v_m\}$, the L2 norm of two images v and u is

$$L2_norm(v, u) = \sqrt{\sum_{i=1}^m (v_i - u_i)^2}$$

Given a large set of unusual images $\{x_1, \dots, x_n\}$, we compute a distance matrix D where $D_{i,j} = L2_norm(x_i, x_j)$. Now we iteratively find unusual images by choosing the image that has the largest distance from all of the images we have chosen so far. Specifically, for each remaining image calculate its distance to our set of exemplars, and choose the image whose distance is the smallest. We define the distance of an image x_0 to a set of exemplars $\{x_1, \dots, x_n\}$ as

$$d = \min_i (D(x_0, x_i))$$



(a)



(b)

Figure 3.3: Figure 3.3(a) shows a montage of interesting images from a golf course webcam. Figure 3.3(b) shows a montage of interesting images of the same scene, but attempts to omit similar images.

In this way, on each iteration, we pick the image that is least likely to be similar to an image we have already selected. We have found that a good way to seed this process is choosing the image with the highest unusualness score as our first exemplar. 3.3(b) shows how similar images are omitted by viewing results in this way.

3.2.3 Two-dimensional Explorer

Using a simple GUI, we can explore a webcam scene in two dimensions. The GUI, shown in ??, displays a a plot of image scores for two different criteria, and displays the image corresponding to a mouseovered datapoint. Using this GUI, it is possible to see which criterias are related and which are not, as well as makes it easier to learn how the criteria effects the scene. [I don't really know what to say in this section].




Figure 3.4: Figure ?? shows a webcam scene displayed using the 2D GUI.

3.3 Criteria

Once webcam scenes are projected onto a PCA basis, there are many different ways to analyze the results. In this section, we will present several different criteria for evaluating this appearance model of a scene, the meaning behind each criteria, and results.

[I don't love the term criteria, but I think it encapsulates the idea of this section...]

3.3.1 PCA coefficient vector magnitude

For each image in a scene, PCA gives a vector of coefficients that correspond to the best linear combination of basis images to reconstruct that image. From this vector, we can assign each image a score equal to the magnitude of this vector. For a vector $v = (v_0, v_1, \dots, v_n)$, the vector magnitude is

$$||v|| = \sqrt{\sum_{i=0}^n v_i^2}$$

This effectively gives us a measure for how far from the mean image in our basis space each image is. Figure 3.5(a) shows the mean image of a webcam scene and Figure 3.5(b) shows several images from that webcam scene that are especially far from the mean.

3.3.2 Residual Error

Given an image and a PCA basis, we can project the image onto the basis and see how well we can reconstruct the image.

[Figure of subplot with image, reconstruction, and residual]


[figure of montage sorted by residualSSD, showing some noisy images and some with real foreground objects]

3.3.3 Variance Model

We can estimate the variance image of a webcam scene as the average of the square of each residual image. This image gives a good summary as to where most activity occurs in a webcam scene. Several examples are shown in Figure ??.



(a)



(b)

Figure 3.5: Figure 3.5(a) shows the mean image of an airport webcam scene. Figure 3.5(b) shows several images found using this technique that are far from the mean.



Figure 3.6: Figure ?? show a webcam scene image, the reconstruction of that image from a PCA basis, and the residual image.



Figure 3.7: Figure ?? show a webcam scene image, the reconstruction of that image from a PCA basis, and the residual image.




Figure 3.8: Figure ?? shows estimated variance images from several webcam scenes.



Figure 3.9: Figure ?? shows several images from the same scene as ??, but with more atypical objects.

Once we have this variance image, we can attempt to isolate independent pixels in a z-score image. We define our z-score image as

$$Z(x, y) = \frac{R(x, y)}{V(x, y)}$$

where $R(x, y)$ is the residual value of the pixel and $V(x, y)$ is the variance value of the pixel. By summing the per pixel z-score for a particular image, we can infer how atypical the residual of that image is. Figure ?? shows several images that contain atypical objects. Such a tool can be used to quickly highlight unusual behavior in a scene.




Figure 3.10: Several residual image histograms and the the best approximation of a normal distribution fitting the data.

3.3.4 Probabilistic

We can also try to learn about an image reconstruction by treating its residual image as samples from an underlying probability density function. If image deviations are due mostly to noise, we predict that these deviations will be normally distributed, but as shown in ??, this is not always the case. We have found that these distributions vary in the same way as webcam scenes, and analyzing them in the right way tells an interesting story.

1. Gaussian Likelihood


The most obvious way to do this is to treat each residual image pixel as a sample from a normal distribution. We can easily estimate the mean and variance of this PDF and then, for each residual value, use the normal distribution equation

$$f(x|\mu, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$


2. Laplacian Likelihood

Many of the residual images have a majority of pixels that are very close zero. This causes the histogram to look very similar to a laplacian distribution.

$$f(x|\mu, \beta) = \frac{1}{2\beta} e^{-\frac{|x-\mu|}{\beta}}$$



(a)



(b)

Figure 3.11: Figure 3.11(a) shows a montage of interesting images from an airport scene that were labeled very non-gaussian. Figure 3.11(b) shows a montage of interesting images from an airport scene that were labeled very non-laplacian.

3. Kurtosis

The kurtosis of a real-valued random variable is a measure of its peakedness. Larger values of kurtosis means more variance is due to less frequent extreme deviations rather than frequent less extreme deviations. It is defined as

$$\gamma_2 = \frac{\mu_4}{\sigma^4}$$

where μ_4 is the fourth moment about the mean and σ^4 is the estimated standard deviation to the fourth power. For a function $f(x)$, the k^{th} moment about the mean is defined as

$$\mu_k = \int_{-\infty}^{\infty} (x - \mu)^k f(x) dx$$

We can use this measure

4. Skewness

The skewness of a random variable is a measure of asymmetry. Higher skewness values mean deviations on one side of the mean do not have corresponding deviations on the other side. It is defined as

$$\gamma_1 = \frac{\mu_3}{\sigma^3}$$

where μ_3 is the third moment about the mean and σ^3 is the cube of the standard deviation.

This measure can be used

Chapter 4

Discussion

The original goal of this project was to find clean mathematical tools to catalog the objects found in a particular webcam scene. Cameras vary in sufficiently difficult ways to make this hard, so we investigated general tools for visualizing these variations.

4.1 Future Work

[Interface with AMOS website] The visualization tools presented in this project

[Item Extraction]

Although the original goal of the project was not met, these visualization tools give information that would be helpful for the original problem.

[Customizable Visualizations]

Another area that this work could be expanded in is user interaction. A tool with which a user could specify what type of information he or she is interested in would be very powerful.

References

- [1] S. Akazawa. The scope of the Japanese information industry in the 1980s. In K. R. Brown, editor, *The challenge of information technology: Proceedings of the forty-first FID (Federation Internationale de Documentation) congress held in Hong Kong 13-16 September 1982*, pages 19–22, Amsterdam, New York, and Oxford, 1983. North Holland.
- [2] American Library Association, Young Adult Services Division, Services Statement Development Committee. Directions for library service to young adults. Chicago, 1978.
- [3] Gilberto Artioli. *Structural studies of the water molecules and hydrogen bonding in zeolites*. PhD thesis, University of Chicago, 1985.
- [4] B. G. F. Cohen. Human aspects in office automation. Technical Report NTIS, PB84-240738, National Institute for Occupational Safety and Health, Division of Biomedical and Behavioral Science, Cincinnati, 1984.
- [5] Martin Greenberger, Julius Aronofsky, James L. McKenney, and William F. Massy, editors. *Networks for research and education: Sharing of computer and information resources nationwide*. MIT Press, Cambridge, 1974.
- [6] Richard Jackson. Running down the up-escalator: Regional inequality in papua new guinea. *Australian Geographer*, 14:175–84, May 1979.
- [7] William Jr. Strunk and E. B. White. *The Elements of Style*. MacMillan Publishing Co., New York, 3rd edition, 1979.
- [8] Kate L. Turabian. *A Manual for Writers of Term Papers, Theses, and Dissertations*. University of Chicago Press, Chicago, 5th edition, 1987.
- [9] University of Chicago Press. *The Chicago Manual of Style*. University of Chicago Press, Chicago, 13th edition, 1982.
- [10] U.S. Department of Agriculture. *Will there be enough food? The 1981 yearbook of agriculture*. Government Printing Office, Washington, D.C., 1981.

Note: Use month and year in which your degree will be conferred.

Tools for Visualizing Webcam Scenes, Ross, M.S. 2009